

# Exercises 6

Brady Lamson

2022-03-28

## 11: Supervised Learning

### 11.1 Classifiers

#### Exercise 1

Based on the decision tree I would make the following predictions:

- A 9-inch, 10-pound pet would likely be a **cat** and be pretty confident.
  - A 14-inch, 21-pound pet would likely be a **dog** and be very confident.
- 

#### Exercise 2

```
type <- c("dog", "dog", "cat", "dog", "cat", "dog", "cat", "dog", "cat",  
         "dog", "cat", "dog", "dog", "cat", "dog", "cat", "dog", "cat", "dog")  
  
wt <- c(8, 17, 8, 18, 7, 22, 6, 16, 7, 20, 10, 15, 14, 11, 13, 13, 15, 17, 10)  
ht <- c(7.5, 10, 8, 15, 7, 15, 7, 13, 11, 16, 7, 10.5, 9, 9.5, 9, 8, 9, 8, 12)  
  
pets <- data.frame(Type = type, Ht = ht, Wt = wt)  
  
my.tree <- rpart(  
  Type ~ Wt + Ht, data = pets,  
  control = rpart.control(minsplit = 7)  
)  
  
newPets <- data.frame(Ht = c(9, 14), Wt = c(10, 21))  
  
predict(my.tree, newdata = newPets, type = "class")  
  
##    1    2  
## cat dog  
## Levels: cat dog
```

First pet is predicted to be a cat and the second pet is predicted to be a dog. This prediction makes the same predictions I made in exercise 1.

---

## Exercise 3

```
my.tree <- rpart(Species ~ Petal.Length + Petal.Width, data = iris)
my.tree

## n= 150
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 150 100 setosa (0.33333333 0.33333333 0.33333333)
##   2) Petal.Length< 2.45 50   0 setosa (1.00000000 0.00000000 0.00000000) *
##   3) Petal.Length>=2.45 100  50 versicolor (0.00000000 0.50000000 0.50000000)
##   6) Petal.Width< 1.75 54   5 versicolor (0.00000000 0.90740741 0.09259259) *
##   7) Petal.Width>=1.75 46   1 virginica (0.00000000 0.02173913 0.97826087) *
```

a) This tree has 3 different terminal nodes.

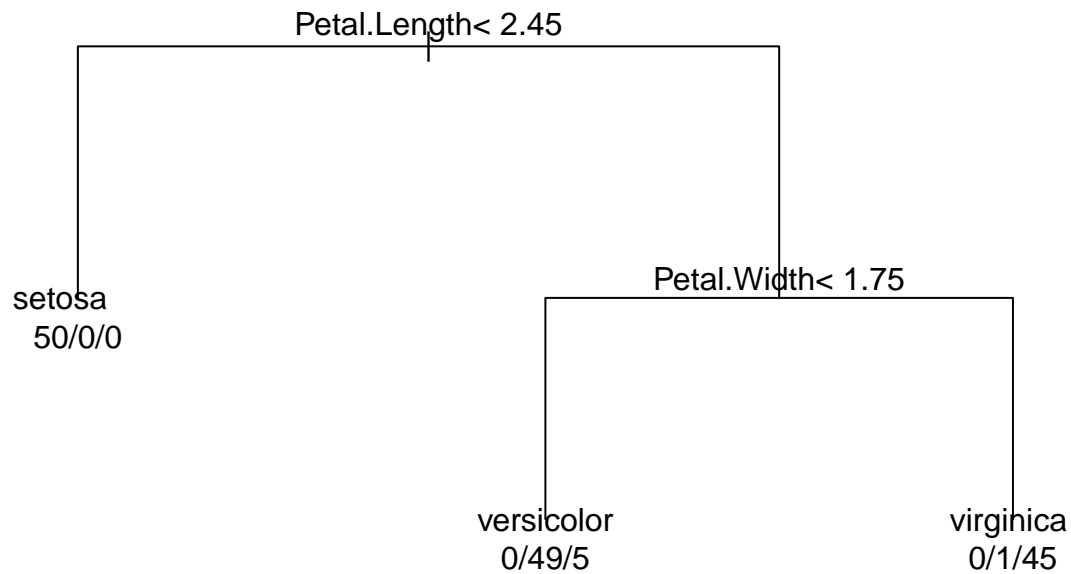
```
predSpecies <- predict(my.tree, type = "class")
species <- data.frame(Actual = iris$Species, Predicted = predSpecies)

confusion <- table(species)
confusion
```

```
##           Predicted
## Actual   setosa versicolor virginica
## setosa      50         0         0
## versicolor   0         49         1
## virginica    0         5         45
```

b) The trees correct classification rate is **96%** and the mis-classification rate is, by extension, **4%**.

```
## First plot:
par(xpd = TRUE)
plot(my.tree, compress = TRUE)
text(my.tree, use.n = TRUE)
```

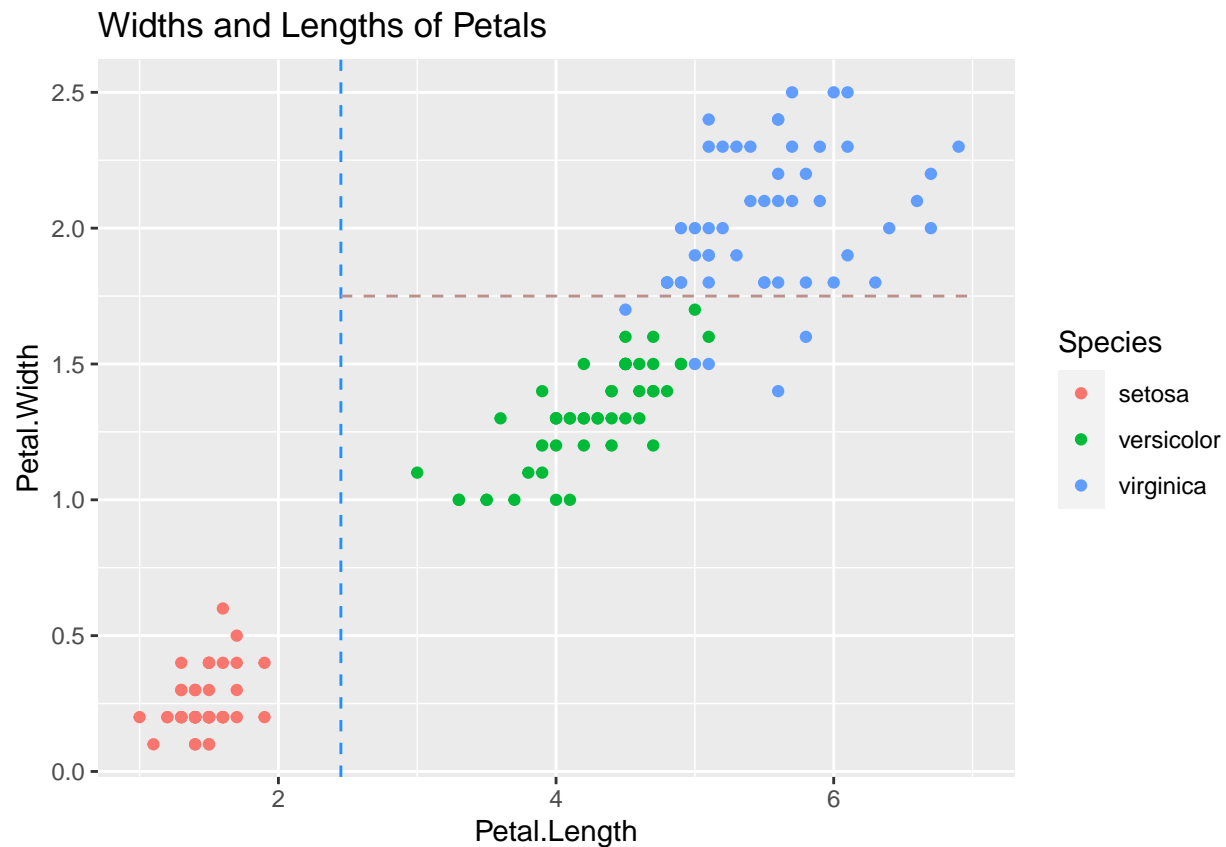


```
par(xpd = FALSE)
```

```
## Second plot:
splitPetal.Length <- 2.45
splitPetal.Width <- 1.75
splitLines <- data.frame(x1 = splitPetal.Length, x2 = 7,
                          y1 = splitPetal.Width, y2 = splitPetal.Width)

g <- ggplot(data = iris,
            mapping = aes(x = Petal.Length, y = Petal.Width,
                          color = Species)) +
  geom_point() +
  labs(title = "Widths and Lengths of Petals") +
  geom_vline(xintercept = splitPetal.Length,
             color = "dodgerblue",
             linetype = 2) +
  geom_segment(data = splitLines,
              mapping = aes(x = x1, y = y1, xend = x2, yend = y2),
              color = "rosybrown", linetype = 2)

g
```



c) Using the above plots I make the following predictions.

- A flower with `Petal.Length` of 3.0cm and `Petal.Width` of 1.5cm would be **versicolor**.

A flower with `Petal.Length` of 4.0cm and `Petal.Width` of 2.1cm would be **virginica**. There aren't any virginica flowers with that length, but the cutoffs on the plot show that's the best prediction to make.

```
newIris <-
  data.frame(Petal.Length = c(3.0, 4.0),
             Petal.Width = c(1.5, 2.1)
  )
```

```
predict(my.tree, newdata = newIris, type = "class")
```

```
##          1          2
## versicolor virginica
## Levels: setosa versicolor virginica
```

The predictions made by the function match that of the predictions I made in part c.

## Exercise 4

```
my.tree <- rpart(Species ~ Sepal.Length + Sepal.Width, data = iris)
my.tree
```

```
## n= 150
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 150 100 setosa (0.33333333 0.33333333 0.33333333)
##    2) Sepal.Length< 5.45 52   7 setosa (0.86538462 0.11538462 0.01923077)
##      4) Sepal.Width>=2.8 45   1 setosa (0.97777778 0.02222222 0.00000000) *
##      5) Sepal.Width< 2.8 7    2 versicolor (0.14285714 0.71428571 0.14285714) *
##    3) Sepal.Length>=5.45 98  49 virginica (0.05102041 0.44897959 0.50000000)
##      6) Sepal.Length< 6.15 43  15 versicolor (0.11627907 0.65116279 0.23255814)
##      12) Sepal.Width>=3.1 7    2 setosa (0.71428571 0.28571429 0.00000000) *
##      13) Sepal.Width< 3.1 36  10 versicolor (0.00000000 0.72222222 0.27777778) *
##      7) Sepal.Length>=6.15 55  16 virginica (0.00000000 0.29090909 0.70909091) *
```

a) This tree has 5 terminal nodes

```
predSpecies <- predict(my.tree, type = "class")
species <- data.frame(Actual = iris$Species, Predicted = predSpecies)
confusion <- table(species)
confusion
```

```
##           Predicted
## Actual   setosa versicolor virginica
## setosa      49         1         0
## versicolor   3        31        16
## virginica    0        11        39
```

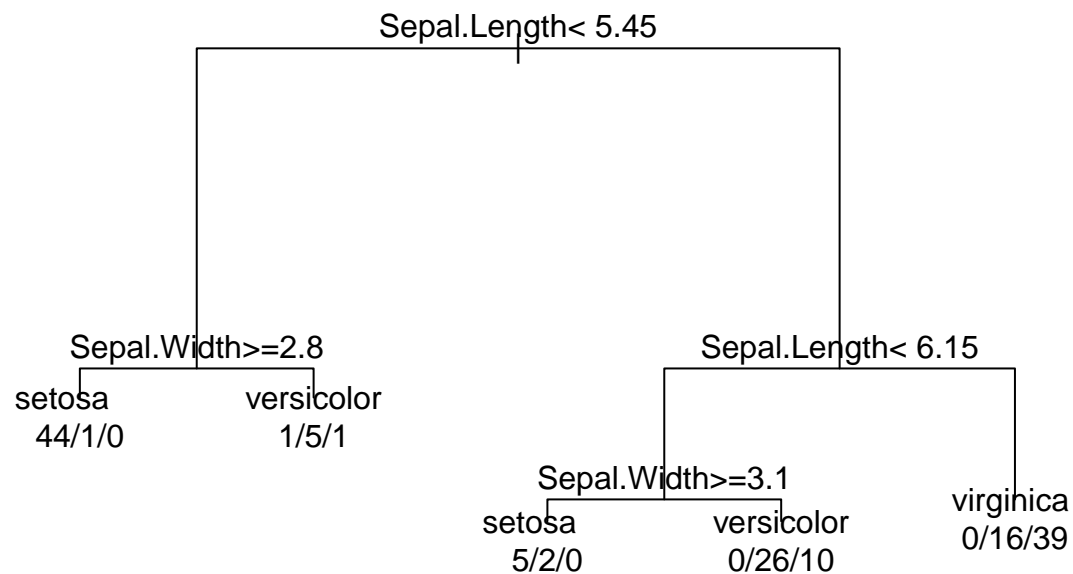
# B

```
classification_rate <- sum(diag(confusion) / nrow(species))

glue::glue("
  The correct classification rate is approximately {(classification_rate * 100) %>% round(3)}%,
  The misclassification rate is approximately {(1 - classification_rate) * 100) %>% round(3)}%
")
```

```
## The correct classification rate is approximately 79.333%,
## The misclassification rate is approximately 20.667%
```

```
## First plot:
par(xpd = TRUE)
plot(my.tree, compress = TRUE)
text(my.tree, use.n = TRUE)
```



```
par(xpd = FALSE)
```

```
## Second plot:
```

```
split1Sepal.Length <- 5.45
split1Sepal.Width <- 2.8
splitLines1 <- data.frame(
  x1 = 3, x2 = split1Sepal.Length,
  y1 = split1Sepal.Width, y2 = split1Sepal.Width
)
```

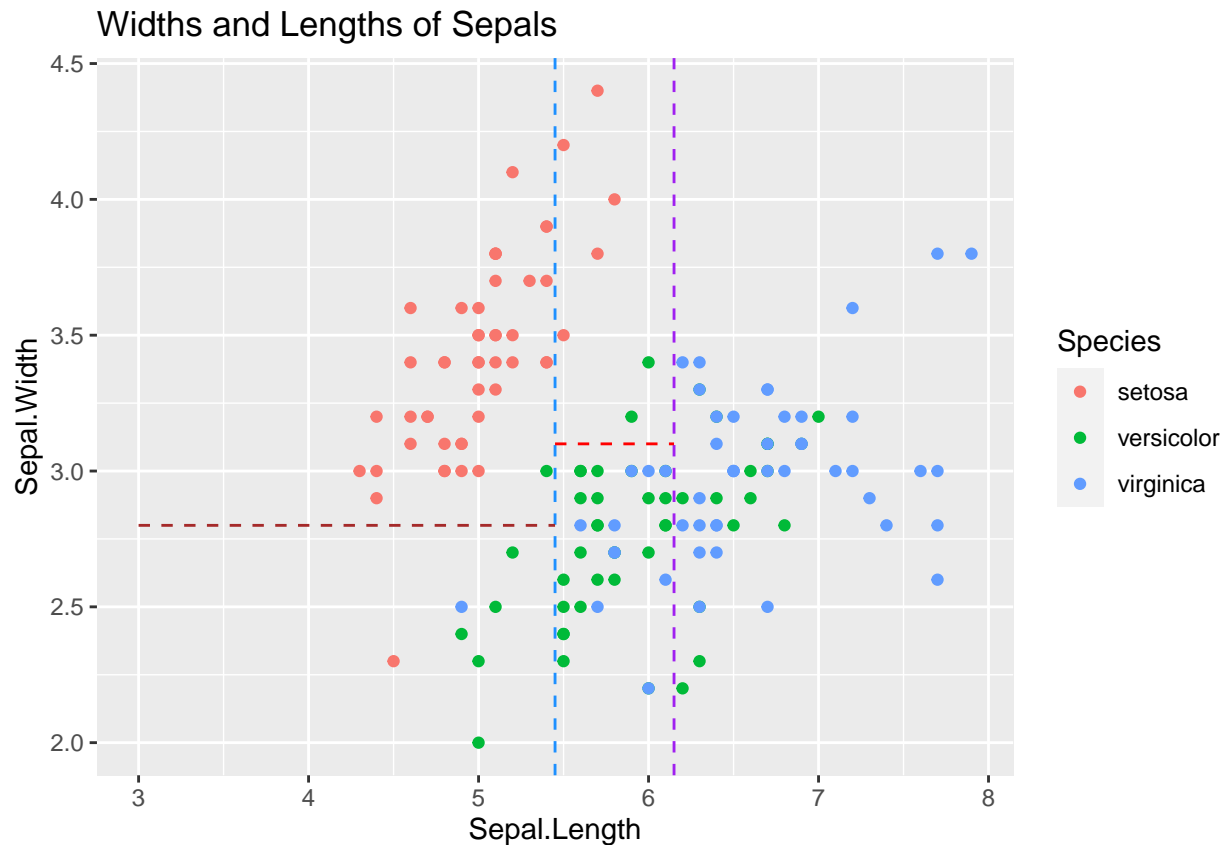
```
split2Sepal.Length <- 6.15
split2Sepal.Width <- 3.1
splitLines2 <- data.frame(
  x1 = split1Sepal.Length, x2 = split2Sepal.Length,
  y1 = split2Sepal.Width, y2 = split2Sepal.Width
)
```

```
g <- ggplot(data = iris,
  mapping = aes(x = Sepal.Length, y = Sepal.Width,
    color = Species)) +
  geom_point() +
  labs(title = "Widths and Lengths of Sepals") +
  geom_vline(xintercept = split1Sepal.Length,
    color = "dodgerblue",
    linetype = 2) +
  geom_vline(xintercept = split2Sepal.Length,
```

```

    color = "purple",
    linetype = 2) +
  geom_segment(data = splitLines1,
    mapping = aes(x = x1, y = y1, xend = x2, yend = y2),
    color = "brown", linetype = 2) +
  geom_segment(data = splitLines2,
    mapping = aes(x = x1, y = y1, xend = x2, yend = y2),
    color = "red", linetype = 2)
g

```



c) Using the above plots I make the following predictions.

- A flower with `Petal.Length` of 6.0cm and `Petal.Width` of 3.5cm would be **setosa**. I used the decision tree over the scatterplot for this decision.

A flower with `Petal.Length` of 7.0cm and `Petal.Width` of 3.0cm would be **virginica**.

```

newIris <- data.frame(Sepal.Length = c(6.0, 7.0),
  Sepal.Width = c(3.5, 3.0))

predict(my.tree, newdata = newIris, type = "class")

```

```

##          1          2
##   setosa virginica
## Levels: setosa versicolor virginica

```

The same decisions as part c were made, the prediction simply appears to match the decision tree cases.

---



## Exercise 5

```
y1 <- c("A", "B", "C", "C", "C", "C", "C", "C", "C", "C", "C", "C")
round(prop.table(table(y1)), digits = 1)
```

```
## y1
##   A   B   C
## 0.1 0.1 0.8
```

```
y2 <- c("A", "B", "C", "C", "A", "C", "B", "A", "A", "B", "C", "B")
round(prop.table(table(y2)), digits = 1)
```

```
## y2
##   A   B   C
## 0.3 0.3 0.3
```

a)

- Based just on the fact that more homogeneous data scores higher on the Gini index, I would guess y1 would be considered more *pure* and thus have a lower score. There are far more C's than anything else. y2 is too evenly split to be considered homogeneous and as such would be further from 0 than y1.

```
1 - sum(
  c(
    .1^2, .1^2, .8^2
  )
)
```

```
## [1] 0.34
```

.34 is the gini index for y1 assuming I did the calculation correctly.

```
1 - sum(
  c(
    .3^2, .3^2, .3^2
  )
)
```

```
## [1] 0.73
```

.73 is the gini index for y2 assuming I did the calculation correctly. This would, at the very least, match my earlier assumptions that this would be higher than the gini index for y1.

- b) The gini index for a vector with all of the same values in them is 0, because the vector is completely pure.  $1 - 1^2 = 0$ .

## Exercise 6

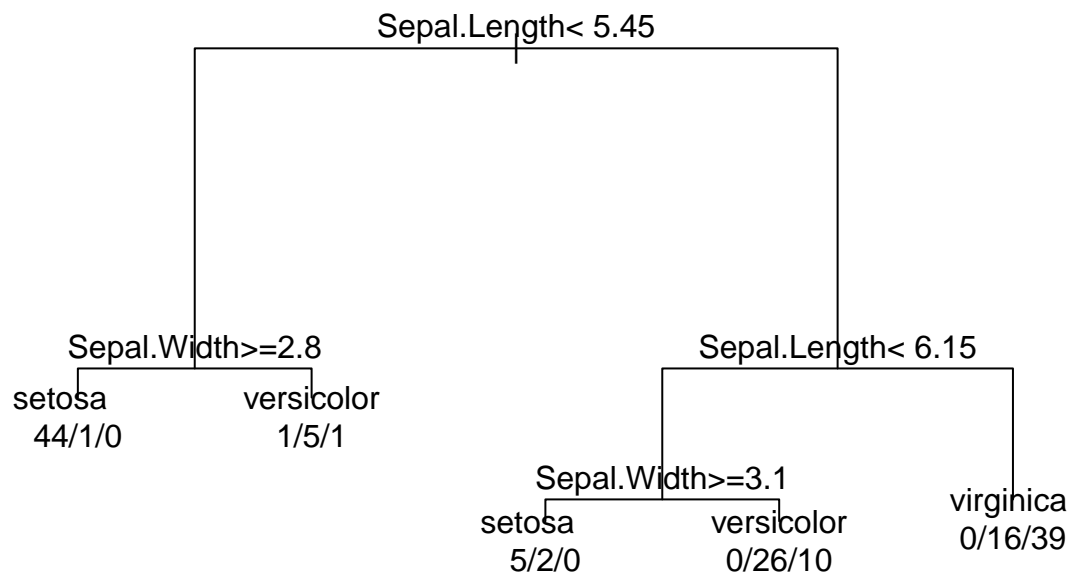
```

my.tree <- rpart(Species ~ Sepal.Length + Sepal.Width, data = iris,
                 control = rpart.control(cp = 0.002))
my.tree

## n= 150
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 150 100 setosa (0.33333333 0.33333333 0.33333333)
##    2) Sepal.Length< 5.45 52   7 setosa (0.86538462 0.11538462 0.01923077)
##      4) Sepal.Width>=2.8 45   1 setosa (0.97777778 0.02222222 0.00000000) *
##      5) Sepal.Width< 2.8 7    2 versicolor (0.14285714 0.71428571 0.14285714) *
##    3) Sepal.Length>=5.45 98  49 virginica (0.05102041 0.44897959 0.50000000)
##      6) Sepal.Length< 6.15 43  15 versicolor (0.11627907 0.65116279 0.23255814)
##        12) Sepal.Width>=3.1 7    2 setosa (0.71428571 0.28571429 0.00000000) *
##        13) Sepal.Width< 3.1 36  10 versicolor (0.00000000 0.72222222 0.27777778) *
##      7) Sepal.Length>=6.15 55  16 virginica (0.00000000 0.29090909 0.70909091) *

par(xpd = TRUE)
plot(my.tree, compress = TRUE)
text(my.tree, use.n = TRUE)

```

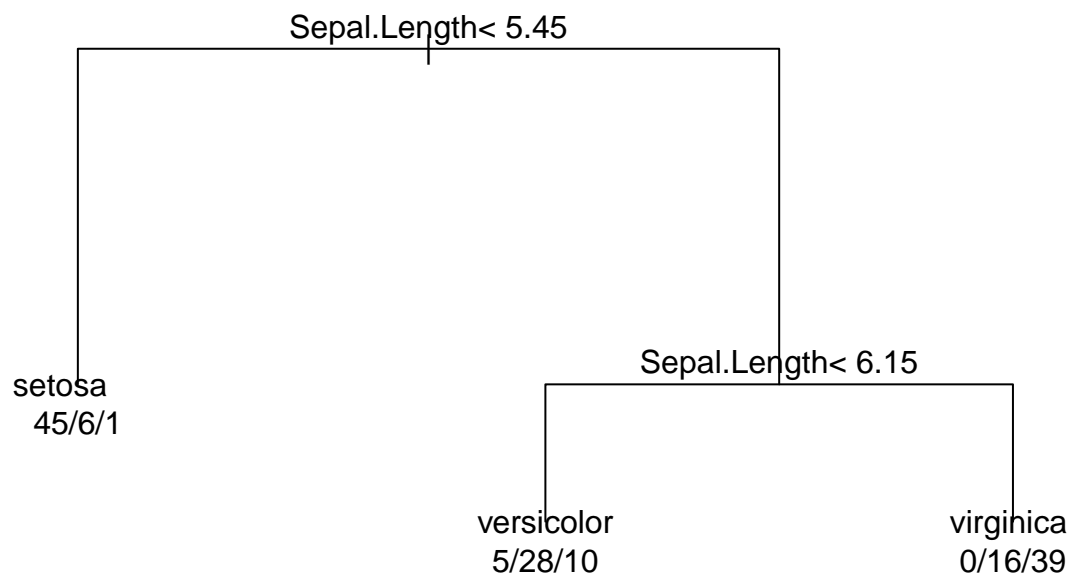


```
par(xpd = FALSE)
```

```
my.tree <- rpart(Species ~ Sepal.Length + Sepal.Width, data = iris,
                 control = rpart.control(cp = 0.05))
my.tree
```

```
## n= 150
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 150 100 setosa (0.33333333 0.33333333 0.33333333)
## 2) Sepal.Length< 5.45 52 7 setosa (0.86538462 0.11538462 0.01923077) *
## 3) Sepal.Length>=5.45 98 49 virginica (0.05102041 0.44897959 0.50000000)
## 6) Sepal.Length< 6.15 43 15 versicolor (0.11627907 0.65116279 0.23255814) *
## 7) Sepal.Length>=6.15 55 16 virginica (0.00000000 0.29090909 0.70909091) *
```

```
par(xpd = TRUE)
plot(my.tree, compress = TRUE)
text(my.tree, use.n = TRUE)
```



```
par(xpd = FALSE)
```

0.2% Resulted in far more terminal nodes.

**Exercise 7**

```
my.forest <- randomForest(Species ~ Sepal.Length + Sepal.Width +  
                          Petal.Length + Petal.Width,  
                          data = iris,  
                          ntree = 500,  
                          mtry = 1)  
  
sum(diag(my.forest$confusion)) / nrow(iris)
```

```
## [1] 0.9466667
```

a) An `mtry` of 1 gives a correct classification rate of approximately 94.67%.

```
my.forest <- randomForest(Species ~ Sepal.Length + Sepal.Width +  
                          Petal.Length + Petal.Width,  
                          data = iris,  
                          ntree = 500,  
                          mtry = 3)  
  
sum(diag(my.forest$confusion)) / nrow(iris)
```

```
## [1] 0.9533333
```

b) An `mtry` of 3 gives a correct classification rate of 96%.

---

## Exercise 8

```
my.forest <- randomForest(Species ~ Sepal.Length + Sepal.Width +
                          Petal.Length + Petal.Width,
                          data = iris,
                          ntree = 500,
                          mtry = 4)

importance(my.forest)
```

```
##              MeanDecreaseGini
## Sepal.Length      1.180067
## Sepal.Width       1.296986
## Petal.Length     44.495274
## Petal.Width      52.378472
```

a) Petal.Width is the most important, Sepal.Length is the least important.

```
newIris <-
  data.frame(Petal.Length = c(3.0, 2.2, 2.7),
             Petal.Width = c(1.2, 2.1, 1.6),
             Sepal.Length = c(5.5, 5.1, 5.9),
             Sepal.Width = c(3.0, 2.7, 3.2))

predict(my.forest, newdata = newIris, type = "class")
```

```
##           1           2           3
## versicolor    setosa versicolor
## Levels: setosa versicolor virginica
```

b) The three predictions were **versicolor**, **setosa**, **versicolor**.

---

## Exercise 9

```
# A

train_iris <- select(iris, -c(Species, Sepal.Length, Sepal.Width))
ret_vec <- c()

for (ktry in 3:8) {

  # k nearest neighbor classification procedure:
  my.knn <- class::knn(train = train_iris,
                      test = train_iris,
                      cl = iris$Species,
                      k = ktry)

  # Save actual and predicted species:
  species <- data.frame(Actual = iris$Species, Predicted = my.knn)

  # Obtain correct classification rate:
  confusion <- table(species)

  ret_vec <- c(ret_vec, sum(diag(confusion)) / nrow(iris))
}

ret_vec
```

```
## [1] 0.9800000 0.9600000 0.9600000 0.9666667 0.9666667 0.9600000
```

```
# B

# Data frame containing new flower for classification:
newIris <- data.frame(Petal.Length = 4.35,
                     Petal.Width = 1.65)
predict_vec <- c()

for (ktry in seq(from = 3, to = 21, by = 3)) {
  # k nearest neighbor classification procedure:
  my.knn <- class::knn(train = train_iris,
                      test = newIris,
                      cl = iris$Species,
                      k = ktry)

  predict_vec <- c(predict_vec, my.knn)
}
predict_vec
```

```
## [1] 2 2 2 2 2 2 2
```

The prediction does not change for any of the provided values of `ktry` I used.

**Exercise 10**

```
rock2 <-
  rock %>%
  dplyr::mutate(
    area = area / 10000,
    peri = peri / 10000,
    perm = log(perm)
  )

ret_vec <- c()

for (ktry in c(1,3,5,7)) {
  my.nn <-
    nnet::nnet(
      perm ~ area + peri + shape, data = rock2,
      size = ktry, linout = TRUE, maxit = 1000, trace = FALSE
    )

  rss <- sum((rock2$perm - predict(my.nn))^2)

  ret_vec <- c(ret_vec, rss)
}

ret_vec
```

```
## [1] 24.052120 10.310556 8.269845 3.806577
```

7 resulted in the smallest residual sum of squares.

---

## Exercise 11

```

# Neural network classification procedure with k = 2 hidden units:
my.nn <-
  nnet::nnet(
    Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
    data = iris, size = 2, maxit = 200, trace = FALSE
  )

newIris <-
  data.frame(
    Petal.Length = c(3.5, 4.7, 1.3),
    Petal.Width = c(1.0, 1.5, 0.5),
    Sepal.Length = c(5.2, 6.2, 4.8),
    Sepal.Width = c(3.3, 3.6, 2.3)
  )

predict(my.nn, newdata = newIris, type = "class")

## [1] "setosa"      "versicolor" "setosa"

print("-----")

## [1] "-----"

predict(my.nn, newdata = newIris)

##      setosa  versicolor   virginica
## 1 1.0000e+00 9.232407e-78 1.428532e-218
## 2 4.9877e-64 1.000000e+00 6.771769e-106
## 3 1.0000e+00 9.232407e-78 1.428532e-218

```

Based on the given probabilities, we can see that **versicolor** has the largest probabilities for the first 2 new flowers. **Setosa** is the highest for the final new flower.



## Exercise 12

```
Ln <- c(85.7, 64.5, 84.1, 82.5, 78.0, 81.3, 71.0, 86.7, 78.7)
Wt <- c(331.9, 121.5, 382.2, 287.3, 224.3, 245.2, 208.2, 393.4, 228.3)
snakes <- data.frame(Length = Ln, Weight = Wt)
```

```
newSnakes <-
  data.frame(
    Length = c(67, 72, 77, 81, 86),
    Weight = c(127.9, 153.7, 204.7, 300.6, 291.4)
  )
```

```
mod0 <- lm(Weight ~ 1, data = snakes)
mod1 <- lm(Weight ~ Length, data = snakes)
mod2 <- lm(Weight ~ poly(Length, 2, raw = TRUE), data = snakes)
mod3 <- lm(Weight ~ poly(Length, 3, raw = TRUE), data = snakes)
mod4 <- lm(Weight ~ poly(Length, 4, raw = TRUE), data = snakes)
mod5 <- lm(Weight ~ poly(Length, 5, raw = TRUE), data = snakes)
```

```
# A
```

```
rmse0 <- sqrt(
  mean(
    (newSnakes$Weight - predict(mod0, newdata = newSnakes))^2
  )
)
```

```
rmse1 <- sqrt(
  mean(
    (newSnakes$Weight - predict(mod1, newdata = newSnakes))^2
  )
)
```

```
rmse2 <- sqrt(
  mean(
    (newSnakes$Weight - predict(mod2, newdata = newSnakes))^2
  )
)
```

```
rmse3 <- sqrt(
  mean(
    (newSnakes$Weight - predict(mod3, newdata = newSnakes))^2
  )
)
```

```
rmse4 <- sqrt(
  mean(
    (newSnakes$Weight - predict(mod4, newdata = newSnakes))^2
  )
)
```

```
rmse5 <- sqrt(
  mean(
```

```

      (newSnakes$Weight - predict(mod5, newdata = newSnakes))^2
    )
  )
rmse0

```

```
## [1] 88.22161
```

```
rmse1
```

```
## [1] 34.56335
```

```
rmse2
```

```
## [1] 36.7128
```

```
rmse3
```

```
## [1] 49.4544
```

```
rmse4
```

```
## [1] 55.22323
```

```
rmse5
```

```
## [1] 53.77739
```

a) The best model according to RMSE is the 2nd. It has the lowest value.

```

mae0 <- mean(abs(newSnakes$Weight - predict(mod0, newdata = newSnakes)))
mae1 <- mean(abs(newSnakes$Weight - predict(mod1, newdata = newSnakes)))
mae2 <- mean(abs(newSnakes$Weight - predict(mod2, newdata = newSnakes)))
mae3 <- mean(abs(newSnakes$Weight - predict(mod3, newdata = newSnakes)))
mae4 <- mean(abs(newSnakes$Weight - predict(mod4, newdata = newSnakes)))
mae5 <- mean(abs(newSnakes$Weight - predict(mod5, newdata = newSnakes)))

mae0

```

```
## [1] 74.96889
```

```
mae1
```

```
## [1] 29.79193
```

```
mae2
```

```
## [1] 29.45119
```

```
mae3
```

```
## [1] 44.99129
```

```
mae4
```

```
## [1] 47.8705
```

```
mae5
```

```
## [1] 45.09594
```

The best model according to MAE is the 3rd model with the 2nd trailing close behind.